

SpringOne TOUR

by VMware Tanzu

Scaling Your Spring Boot App to Zero

DaShaun Carter
Spring Developer Advocate, VMware

Thomas Risberg
Staff Engineer, VMware



Tweet



Kelsey Hightower

@kelseyhightower



Scale to zero is a trade-off; not a best practice.

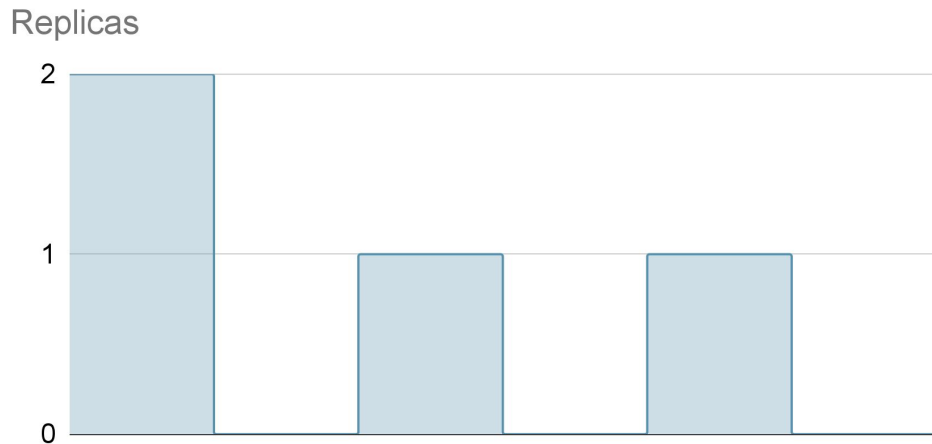
11:58 AM · Apr 16, 2019

37 Retweets **4** Quotes **226** Likes **2** Bookmarks

<https://twitter.com/kelseyhightower/status/1118181786914136064>

Scale to zero

- Generalization of serverless
- Typically hosted on Kubernetes platforms based on Knative/KEDA
- Not limited to functions
- Pay-as-you-use billing model



**2 ways to scale your
Spring applications to zero
with sub-second startup**

1

Scale to zero with native Spring Boot and GraalVM

Spring AOT GraalVM Native Images

A new deployment variant


As an alternative to classic HotSpot deployment

```
$ ./gradlew bootBuildImage
Successfully built image 'docker.io/library/spring-petclinic:1.0.0-SNAPSHOT'
```

```
$ docker run --rm -p 8080:8080 spring-petclinic:1.0.0-SNAPSHOT
```

[illegible]

```
:: Built with Spring Boot :: 3.1.4
```

```
Starting AOT-processed PetClinicApplication using Java 17.0.8.1 with PID 1
No active profile set, falling back to 1 default profile: "default"
Tomcat initialized with port(s): 8080 (http)
Starting service [Tomcat]
Starting Servlet engine: [Apache Tomcat/10.1.13]
Initializing Spring embedded WebApplicationContext
Root WebApplicationContext: initialization completed in 19 ms
HikariPool-1 - Starting...
HikariPool-1 - Added connection conn0: url=jdbc:h2:mem:testdb user=PETCLINIC
HikariPool-1 - Start completed.
Exposing 13 endpoint(s) beneath base path '/actuator'
Tomcat started on port(s): 8080 (http) with context path ''
Started PetClinicApplication in 0.107 seconds (process running for 0.109) 
```

Sub-second startup on production!

GraalVM advantages



Instant startup

Milliseconds for native instead of seconds for the JVM



No warmup

Peak performance available immediately



Low resource usage

Lower memory footprint and no JIT compilation



Reduced surface attack

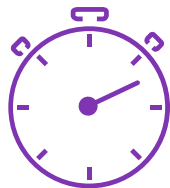
Closed world of dependencies with explicit reflection and serialization



Compact packaging

Smaller container easier to deploy

GraalVM trade-offs



Very slow compilation

Minutes instead of seconds



Compatibility

Reachability metadata required for reflection, proxies, resources



Closed-world Assumptions

Bean conditions fixed at build time
No dynamic class loading



Thomas Wuerthinger

@thomaswue

“The startup time of the application was reduced from approximately 30 seconds down to about 3 ms, and more importantly the memory usage was also significantly reduced from 6.6 GB down to 1 GB, with the same throughput and CPU utilization.” 👍🚀 @graalvm



Thomas Schuehly @tschuehly · 21 Jul

10minutemail.com runs on @GraalVM @springboot native image with server-side rendering with @thymeleaf 🤖
digitalsanctuary.com/10minutemail/m...

<https://twitter.com/thomaswue/status/1682465475748298755>



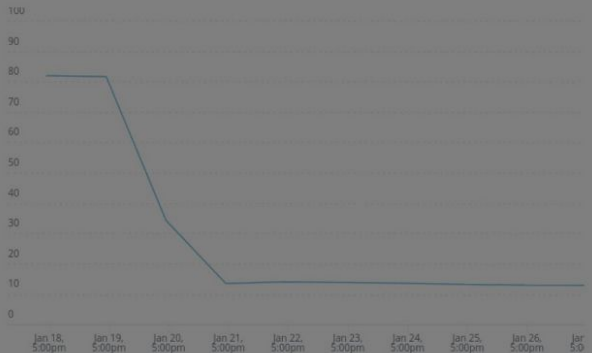
Thomas Schuehly

@tschuehly

10minutemail.com runs on @GraalVM @springboot native image with server-side rendering with @thymeleaf 🤖
digitalsanctuary.com/10minutemail/m...

Memory usage (%)

Since Jan 18, 04:00 pm until Jan 28, 04:00 pm



Migrating 10MinuteMail from Java to GraalVM Native

12:59 pm · 21 Jul 2023 · 36K Views

16 Retweets · 2 Quotes · 88 Likes · 29 Bookmarks



Thomas Wuerthinger

@thomaswue

Big news from last week's @GraalVM for @Java 21 release was that we can finally demonstrate that AOT (with profile-guided optimizations) can outperform JIT for all major metrics including throughput! Here are the numbers for the @Springframework PetClinic example (both G1 GC). 🚀

	GraalVM CE with C2 JIT	Oracle GraalVM Native Image	
Memory Usage (max RSS)	1,029 MB	641 MB	-38% lower
Peak throughput	11,066 req/s	11,902 req/s	+8% higher
Throughput per memory	12,488 req/(GB*s)	18,569 req/(GB*s)	+49% better
Tail latency (P99)	7.2ms	5.15ms	-28% lower
Startup	7,090ms	210ms	34x faster

11:44 AM · Sep 29, 2023 · 25.3K Views

<https://twitter.com/thomaswue/status/1707783370187440148>



Demo - *Spring PetClinic native build running as Knative service*



2

Scale to zero with the JVM Spring Boot and Project CRaC

JVM Checkpoint Restore

Project CRaC

Immediate startup for Spring deployments on HotSpot



Coming up...

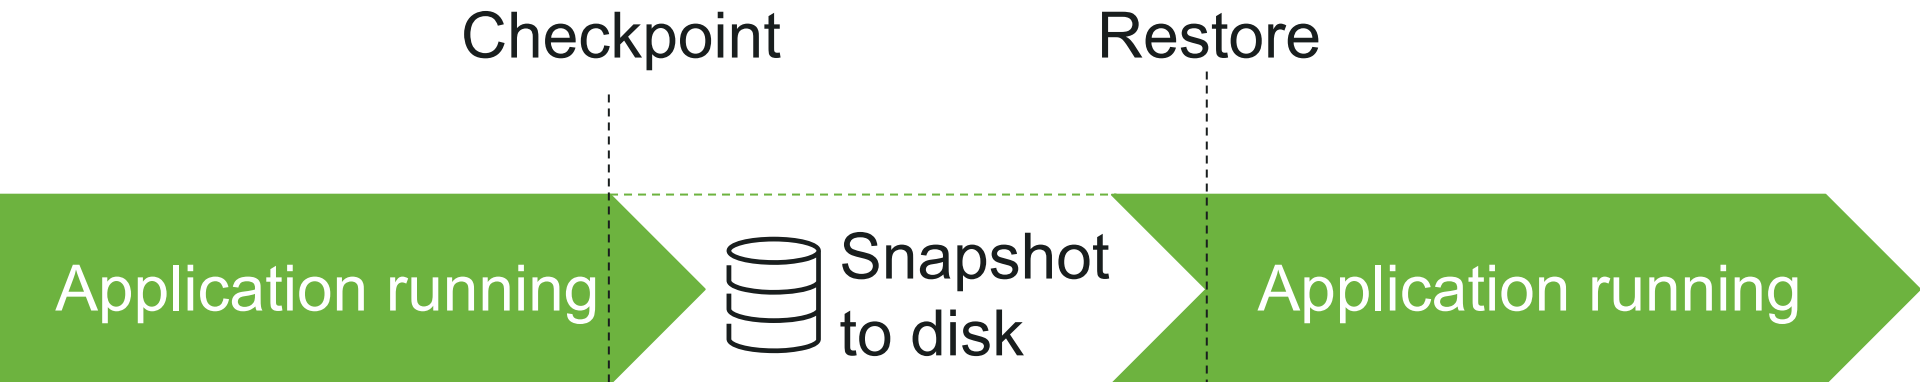


Spring Framework 6.1

Spring Boot 3.2

November 2023

JVM Checkpoint Restore



**Spring Boot 3.2 introduces
initial support
for JVM Checkpoint Restore
with Project CRaC**

Check the evolving scope of CRaC support

<https://github.com/spring-projects/spring-checkpoint-restore-smoke-tests/blob/main/STATUS.adoc>

Boot

Smoke Test	appTest	checkpointRestoreAppTest	test
actuator-webflux	build passing	build passing	
actuator-webmvc	build passing	build passing	

Cloud

Smoke Test	appTest	checkpointRestoreAppTest	test
context-refresh	build passing	build passing	build passing
context-refresh-http	build passing	build passing	build passing

Data

Smoke Test	appTest	checkpointRestoreAppTest	test
data-jdbc	build passing	build passing	
data-jpa	build passing	build passing	
data-redis	build passing	build passing	

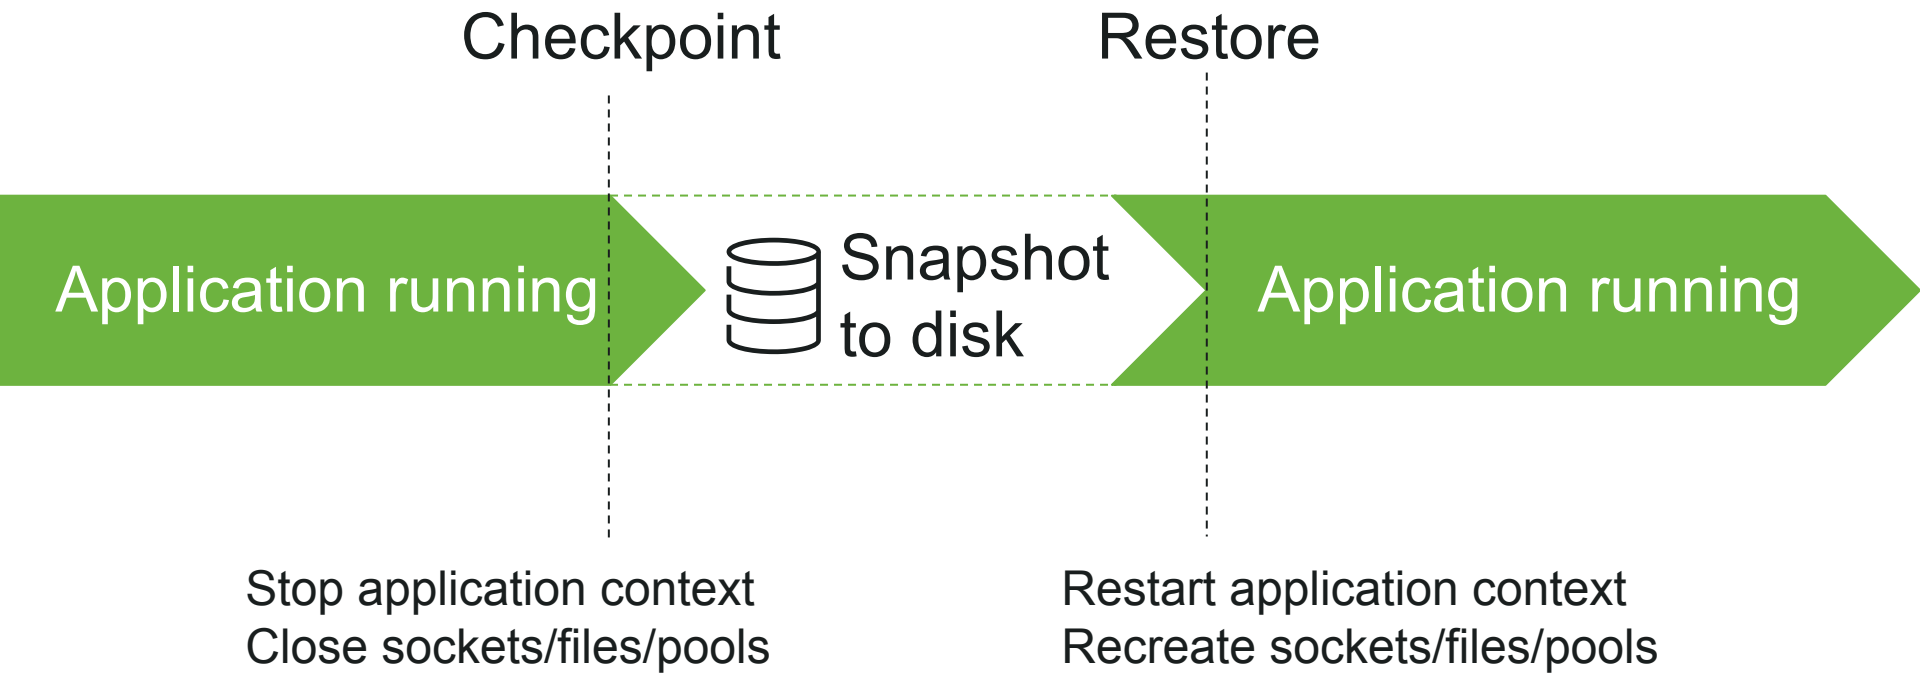
Framework

Smoke Test	appTest	checkpointRestoreAppTest	test
hibernate-mysql	build passing	build passing	
scheduled	build passing	build passing	
webclient-netty	build passing	build passing	
webflux-netty	build passing	build passing	build passing
webflux-undertow	build passing	build passing	build passing
webmvc-jetty	build passing	build passing	build passing
webmvc-tomcat	build passing	build passing	build passing

Integration

Smoke Test	appTest	checkpointRestoreAppTest	test
integration-basic	build passing	build passing	
integration-webflux-data	build passing	build passing	
spring-amqp-rabbit	build passing	build passing	
spring-kafka	build passing	build passing	
spring-kafka-avro	build passing	build passing	
spring-kafka-streams	build passing	build passing	

Spring lifecycle with CRaC



```
:: Built with Spring Boot :: 3.2.0-M3
```

Optional load testing
to warm up the JVM
then checkpoint

Sub-second restoration on production with the JVM!

Regular startup
typically on
CI/CD platform

Project CRaC advantages



Instant startup

Milliseconds
instead of seconds
for the regular JVM



No runtime warmup

Peak performance
available immediately



Compatibility

Still a JVM with all
regular capabilities

Project CRaC trade-offs



Checkpoint startup

Require to start the application ahead



Lifecycle management

Require to close and reopen sockets, files, pools



Secret management

Sensitive informations may leak in the snapshot files



System Integration

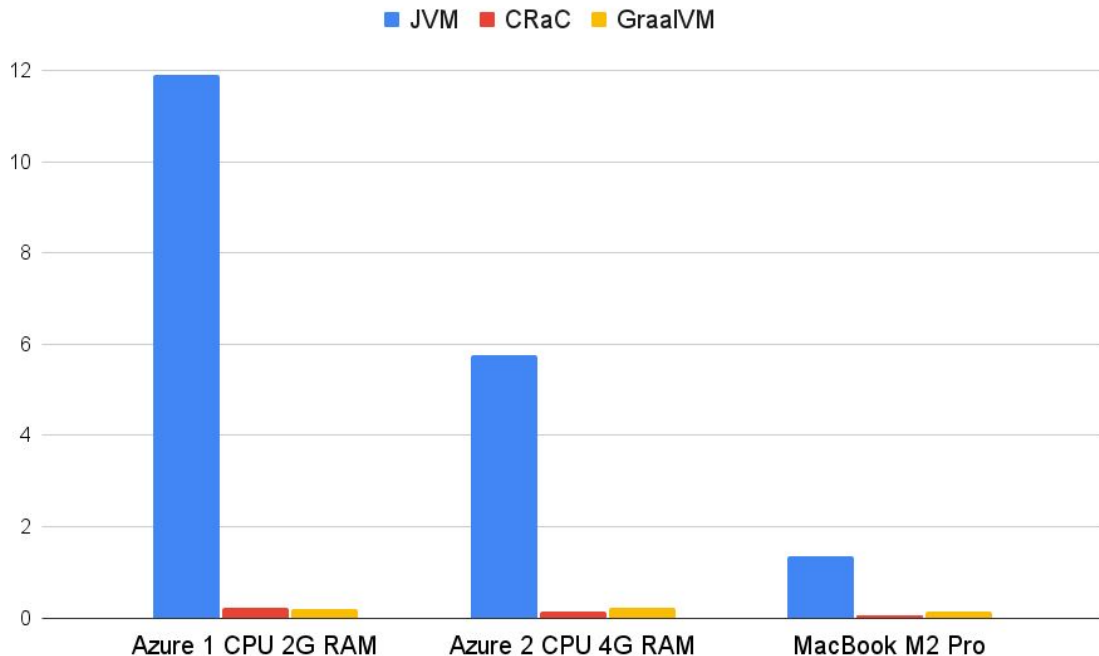
Linux only and advanced capability management required

Demo - *Getting started with JVM Checkpoint/Restore builds*



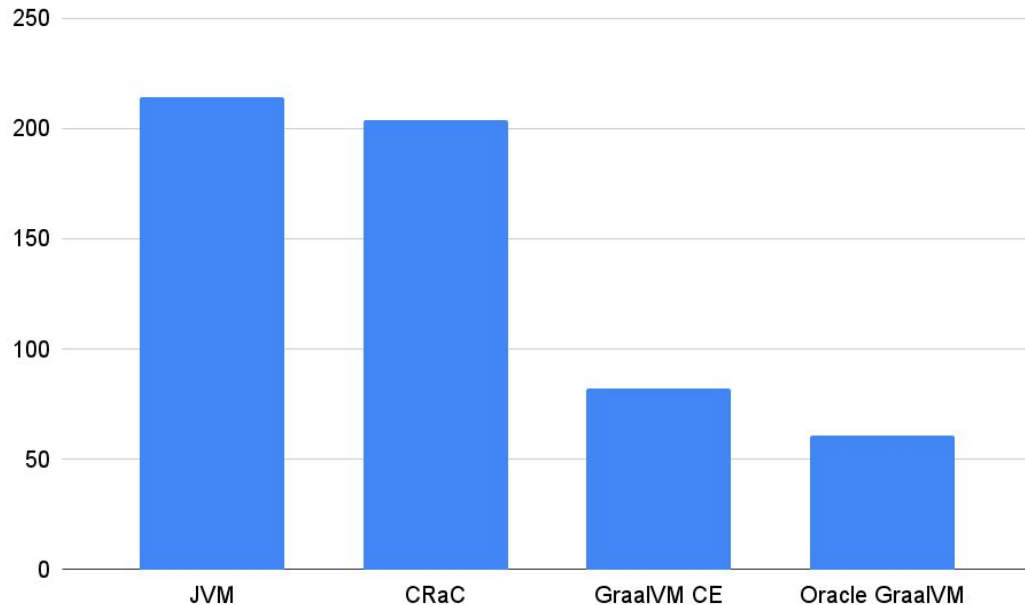
Petclinic startup time (seconds)

**GraalVM and CRaC
both allow 50x faster
startup time and
scaling to zero**

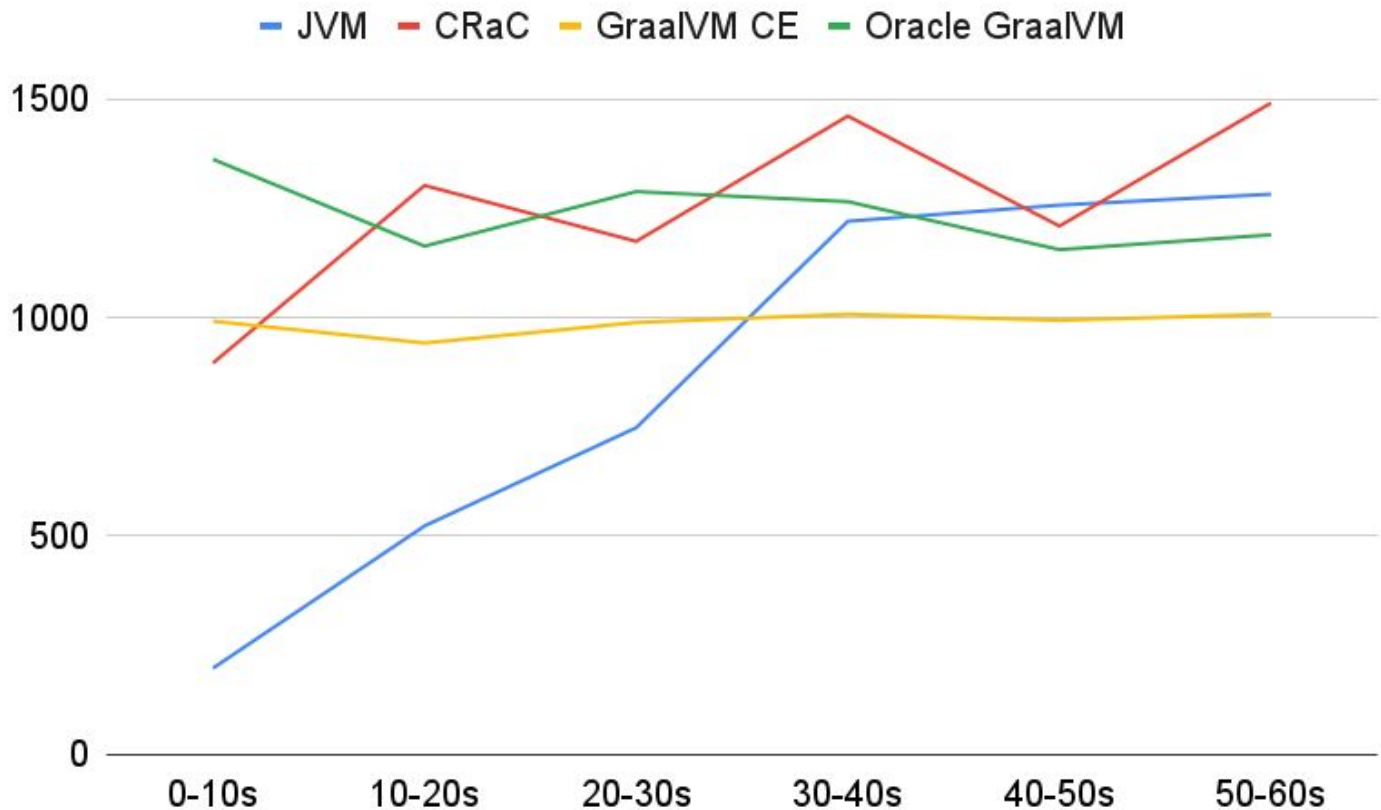


Petclinic memory consumption after startup (MB)

GraalVM allows 3x reduction of memory consumption of Spring Boot application infrastructure



Petclinic throughput on 1 CPU 2G RAM (req/s)



Thank you

Follow us

@dashaun

@trisberg

Dig Deeper:

<https://github.com/trisberg/springone-tour-2023>

SpringOne TOUR

by VMware Tanzu